

# NLP Assignment 1 (40% of grade): Text classification for Fake News Detection

This coursework will involve you implementing functions for a text classifier, which you will train to detect **fake news** in a corpus of approx. 10,000 statements, which will be split into a 80%/20% training/test split.

In this template you are given the basis for that implementation, though some of the functions are missing, which you have to fill in.

Follow the instructions file **NLP\_Assignment\_1\_Instructions.pdf** for details of each question - the outline of what needs to be achieved for each question is as below.

You must submit all **ipython notebooks and extra resources you need to run the code if you've added them** in the code submission, and a **2 page report (pdf)** in the report submission on QMPlus where you report your methods and findings according to the instructions file for each question.

```
In [1]: !pip install nltk
```

```
Requirement already satisfied: nltk in /opt/conda/lib/python3.10/site-packages (3.7)
Requirement already satisfied: tqdm in /opt/conda/lib/python3.10/site-packages (from nltk) (4.64.1)
Requirement already satisfied: click in /opt/conda/lib/python3.10/site-packages (from nltk) (8.1.3)
Requirement already satisfied: regex<=2021.8.3 in /opt/conda/lib/python3.10/site-packages (from nltk) (2022.10.31)
Requirement already satisfied: joblib in /opt/conda/lib/python3.10/site-packages (from nltk) (1.2.0)
```

```
[notice] A new release of pip available: 22.3 -> 22.3.1
```

```
[notice] To update, run: pip install --upgrade pip
```

```
In [2]: import pandas as pd
import csv # csv reader
from sklearn.svm import LinearSVC
from nltk.classify import SklearnClassifier
from sklearn.pipeline import Pipeline
from sklearn.metrics import precision_recall_fscore_support # to report o
import numpy as np
import nltk
from nltk.corpus import stopwords
from nltk.stem import PorterStemmer
from nltk.classify import accuracy
```

## Question 1: Input and Basic preprocessing (10 marks)

```
In [3]: def convert_label(label):
        """Converts the multiple classes into two,
        making it a binary distinction between fake news and real."""
        #return label
        # Converting the multiclass labels to binary label
        labels_map = {
            'true': 'REAL',
            'mostly-true': 'REAL',
            'half-true': 'REAL',
            'false': 'FAKE',
            'barely-true': 'FAKE',
            'pants-fire': 'FAKE'
        }
        return labels_map[label]

def parse_data_line(data_line):
    # Should return a tuple of the label as just FAKE or REAL and the sta
    # e.g. (label, statement)
    data_line[0]=convert_label(data_line[0])
    a=(data_line[0],data_line[1])
    return a
```

```
In [4]: # Input: a string of one statement
def pre_process(text):
    # Should return a list of tokens
    # DESCRIBE YOUR METHOD IN WORDS
    a=nlTK.word_tokenize(text)
    interpuations = [',', '.', ':', ';', '?', '(', ')', '[', ']', '&'
    cutwords2 = [word for word in a if word not in interpuations]
    stops = set(stopwords.words("english"))
    cutwords3 = [word for word in cutwords2 if word not in stops]
    cutwords4 = []
    for cutword in cutwords3:
        cutwords4.append(PorterStemmer().stem(cutword))
    return cutwords4
```

## Question 2: Basic Feature Extraction (20 marks)

```
In [5]: def to_feature_vector(tokens):
        global_feature_dict = {}
        # Should return a dictionary containing features as keys, and weights
        # DESCRIBE YOUR METHOD IN WORDS
        for item in tokens:
            global_feature_dict[item]=global_feature_dict.get(item,0)+1
        return global_feature_dict
```

```
In [6]: def load_data(path):
        """Load data from a tab-separated file and append it to raw_data."""
        raw_data=[]
        reader = pd.read_csv(path,delimiter='\t')
        for i in range(len(reader)):
            (label, text) = parse_data_line(reader.iloc[i,:])
            raw_data.append((text, label))
        return raw_data

        def split_and_preprocess_data(percentage,path):
            """Split the data between train_data and test_data according to the p
            and performs the preprocessing."""
            train_data=[]
            test_data=[]
            raw_data=load_data(path)
            num_samples = len(raw_data)
            num_training_samples = int((percentage * num_samples))
            for (text, label) in raw_data[:num_training_samples]:
                train_data.append((to_feature_vector(pre_process(text)),label))
            for (text, label) in raw_data[num_training_samples:]:
                test_data.append((to_feature_vector(pre_process(text)),label))
            return train_data ,test_data
```

```
In [7]: # TRAINING AND VALIDATING OUR CLASSIFIER
```

```
def train_classifier(data):
    print("Training Classifier...")
    pipeline = Pipeline([('svc', LinearSVC())])
    model=SklearnClassifier(pipeline).train(data)
    return model
```

```
In [8]: train_data ,test_data=split_and_preprocess_data(0.9,'fake_news.tsv')
```

```
/tmp/ipykernel_258/909790341.py:20: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame

See the caveats in the documentation: https://pandas.pydata.org/pandas-d
ocs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy
    data_line[0]=convert_label(data_line[0])

-----
LookupError                                Traceback (most recent call la
st)
Cell In [8], line 1
```

```

----> 1 train_data ,test_data=split_and_preprocess_data(0.9,'fake_news.t
sv')

Cell In [6], line 19, in split_and_preprocess_data(percentage, path)
    17 num_training_samples = int((percentage * num_samples))
    18 for (text, label) in raw_data[:num_training_samples]:
----> 19     train_data.append((to_feature_vector(pre_process(text)),label))
    20 for (text, label) in raw_data[num_training_samples:]:
    21     test_data.append((to_feature_vector(pre_process(text)),label))

Cell In [4], line 5, in pre_process(text)
    2 def pre_process(text):
    3     # Should return a list of tokens
    4     # DESCRIBE YOUR METHOD IN WORDS
----> 5     a=nlk.word_tokenize(text)
    6     interpuations = [',', '.', ':', ';', '?', '(', ')', '[',
    7     ']', '&', '!', '*', '@', '#', '$', '%']
    7     cutwords2 = [word for word in a if word not in
interpuations]

File /opt/conda/lib/python3.10/site-packages/nltk/tokenize/__init__.py:1
29, in word_tokenize(text, language, preserve_line)
    114 def word_tokenize(text, language="english", preserve_line=False)
:
    115     """
    116     Return a tokenized copy of *text*,
    117     using NLTK's recommended word tokenizer
    (...)
    127     :type preserve_line: bool
    128     """
--> 129     sentences = [text] if preserve_line else sent_tokenize(text,
language)
    130     return [
    131         token for sent in sentences for token in
_treebank_word_tokenizer.tokenize(sent)
    132     ]

File /opt/conda/lib/python3.10/site-packages/nltk/tokenize/__init__.py:1
06, in sent_tokenize(text, language)
    96 def sent_tokenize(text, language="english"):
    97     """
    98     Return a sentence-tokenized copy of *text*,
    99     using NLTK's recommended sentence tokenizer
    (...)
    104     :param language: the model name in the Punkt corpus
    105     """
--> 106     tokenizer = load(f"tokenizers/punkt/{language}.pickle")
    107     return tokenizer.tokenize(text)

```

```

File /opt/conda/lib/python3.10/site-packages/nltk/data.py:750, in load(r
esource_url, format, cache, verbose, logic_parser, fstruct_reader, encod

```

```

ing)
747     print(f"<<Loading {resource_url}>>")
749 # Load the resource.
--> 750 opened_resource = open(resource_url)
752 if format == "raw":
753     resource_val = opened_resource.read()

File /opt/conda/lib/python3.10/site-packages/nltk/data.py:876, in _open(
resource_url)
873 protocol, path_ = split_resource_url(resource_url)
875 if protocol is None or protocol.lower() == "nltk":
--> 876     return find(path_, path + [""]).open()
877 elif protocol.lower() == "file":
878     # urllib might not use mode='rb', so handle this one ourselv
es:
879     return find(path_, [""]).open()

File /opt/conda/lib/python3.10/site-packages/nltk/data.py:583, in find(r
esource_name, paths)
581 sep = "*" * 70
582 resource_not_found = f"\n{sep}\n{msg}\n{sep}\n"
--> 583 raise LookupError(resource_not_found)

LookupError:
*****
Resource punkt not found.
Please use the NLTK Downloader to obtain the resource:

>>> import nltk
>>> nltk.download('punkt')

For more information see: https://www.nltk.org/data.html

Attempted to load tokenizers/punkt/PY3/english.pickle

Searched in:
- '/home/jovyan/nltk_data'
- '/opt/conda/nltk_data'
- '/opt/conda/share/nltk_data'
- '/opt/conda/lib/nltk_data'
- '/usr/share/nltk_data'
- '/usr/local/share/nltk_data'
- '/usr/lib/nltk_data'
- '/usr/local/lib/nltk_data'
- ''
*****

```

## Question 3: Cross-validation (20 marks)

```
In [9]: #solution
from sklearn.metrics import classification_report

def cross_validate(dataset, folds):
    results = []
    fold_size = int(len(dataset)/folds) + 1

    for i in range(0, len(dataset), int(fold_size)):
        # insert code here that trains and tests on the 10 folds of data
        result=train_classifier(dataset[i:i+fold_size])
        results.append(accuracy(result, test_data))
        print("Fold start on items %d - %d" % (i, i+fold_size))
        # FILL IN THE METHOD HERE

    return results
```

```
In [10]: # PREDICTING LABELS GIVEN A CLASSIFIER

def predict_labels(samples, classifier):
    """Assuming preprocessed samples, return their predicted labels from
    return classifier.classify_many(samples)

def predict_label_from_raw(sample, classifier):
    """Assuming raw text, return its predicted label from the classifier
    return classifier.classify(to_feature_vector(pre_process(sample)))
```

In [11]: # MAIN

```

# loading reviews
# initialize global lists that will be appended to by the methods below
raw_data = []          # the filtered data from the dataset file
train_data = []        # the pre-processed training data as a percentage
test_data = []         # the pre-processed test data as a percentage of t

# references to the data files
data_file_path = 'fake_news.tsv'

# Do the actual stuff (i.e. call the functions we've made)
# We parse the dataset and put it in a raw data list
print("Now %d rawData, %d trainData, %d testData" % (len(raw_data), len(t
    "Preparing the dataset...",sep='\n')

raw_data=load_data(data_file_path)

# We split the raw dataset into a set of training data and a set of test
# You do the cross validation on the 80% (training data)
# We print the number of training samples and the number of features befo
print("Now %d rawData, %d trainData, %d testData" % (len(raw_data), len(t
    "Preparing training and test data...",sep='\n')

train_data,test_data=split_and_preprocess_data(0.8,data_file_path)
global_feature_dict=to_feature_vector(train_data[0][0])

# We print the number of training samples and the number of features afte
print("After split, %d rawData, %d trainData, %d testData" % (len(raw_dat
    "Training Samples: ", len(train_data), "Features: ", len(global_fea

```

Now 0 rawData, 0 trainData, 0 testData

Preparing the dataset...

/tmp/ipykernel\_258/909790341.py:20: SettingWithCopyWarning:  
A value is trying to be set on a copy of a slice from a DataFrame

See the caveats in the documentation: [https://pandas.pydata.org/pandas-docs/stable/user\\_guide/indexing.html#returning-a-view-versus-a-copy](https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy)  
data\_line[0]=convert\_label(data\_line[0])

Now 10241 rawData, 0 trainData, 0 testData

Preparing training and test data...

/tmp/ipykernel\_258/909790341.py:20: SettingWithCopyWarning:  
A value is trying to be set on a copy of a slice from a DataFrame

See the caveats in the documentation: [https://pandas.pydata.org/pandas-docs/stable/user\\_guide/indexing.html#returning-a-view-versus-a-copy](https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy)  
data\_line[0]=convert\_label(data\_line[0])

```

-----
LookupError                                Traceback (most recent call la
st)

```

```

Cell In [11], line 26
    20 # We split the raw dataset into a set of training data and a set
of test data (80/20)
    21 # You do the cross validation on the 80% (training data)
    22 # We print the number of training samples and the number of feat
ures before the split
    23 print("Now %d rawData, %d trainData, %d testData" % (len(raw_dat
a), len(train_data), len(test_data)),
    24         "Preparing training and test data...",sep='\n')
--> 26 train_data,test_data=split_and_preprocess_data(0.8,data_file_pat
h)
    27 global_feature_dict=to_feature_vector(train_data[0][0])
    29 # We print the number of training samples and the number of feat
ures after the split

```

```

Cell In [6], line 19, in split_and_preprocess_data(percentage, path)
    17 num_training_samples = int((percentage * num_samples))
    18 for (text, label) in raw_data[:num_training_samples]:
--> 19     train_data.append((to_feature_vector(pre_process(text)),label))
    20 for (text, label) in raw_data[num_training_samples:]:
    21     test_data.append((to_feature_vector(pre_process(text)),label))

```

```

Cell In [4], line 5, in pre_process(text)
    2 def pre_process(text):
    3     # Should return a list of tokens
    4     # DESCRIBE YOUR METHOD IN WORDS
----> 5     a=nlk.word_tokenize(text)
    6     interpuations = [',', '.', ':', ';', '?', '(', ')', '[',
',', '&', '!', '*', '@', '#', '$', '%']
    7     cutwords2 = [word for word in a if word not in
interpuations]

```

```

File /opt/conda/lib/python3.10/site-packages/nltk/tokenize/__init__.py:1
29, in word_tokenize(text, language, preserve_line)
    114 def word_tokenize(text, language="english", preserve_line=False)
:
    115     """
    116     Return a tokenized copy of *text*,
    117     using NLTK's recommended word tokenizer
    (...)
    127     :type preserve_line: bool
    128     """
--> 129     sentences = [text] if preserve_line else sent_tokenize(text,
language)
    130     return [
    131         token for sent in sentences for token in
_treebank_word_tokenizer.tokenize(sent)
    132     ]

```

```

File /opt/conda/lib/python3.10/site-packages/nltk/tokenize/__init__.py:1
06, in sent_tokenize(text, language)

```



```

96 def sent_tokenize(text, language="english"):
97     """
98     Return a sentence-tokenized copy of *text*,
99     using NLTK's recommended sentence tokenizer
100     (...)
101     :param language: the model name in the Punkt corpus
102     """
--> 106 tokenizer = load(f"tokenizers/punkt/{language}.pickle")
107 return tokenizer.tokenize(text)

```

File /opt/conda/lib/python3.10/site-packages/nltk/data.py:750, in load(resource\_url, format, cache, verbose, logic\_parser, fstruct\_reader, encoding)

```

747     print(f"<<Loading {resource_url}>>")
749 # Load the resource.
--> 750 opened_resource = _open(resource_url)
752 if format == "raw":
753     resource_val = opened_resource.read()

```

File /opt/conda/lib/python3.10/site-packages/nltk/data.py:876, in \_open(resource\_url)

```

873 protocol, path_ = split_resource_url(resource_url)
875 if protocol is None or protocol.lower() == "nltk":
--> 876     return find(path_, path + []).open()
877 elif protocol.lower() == "file":
878     # urllib might not use mode='rb', so handle this one ourselves:
879     return find(path_, []).open()

```

File /opt/conda/lib/python3.10/site-packages/nltk/data.py:583, in find(resource\_name, paths)

```

581 sep = "*" * 70
582 resource_not_found = f"\n{sep}\n{msg}\n{sep}\n"
--> 583 raise LookupError(resource_not_found)

```

LookupError:

\*\*\*\*\*

Resource punkt not found.

Please use the NLTK Downloader to obtain the resource:

```

>>> import nltk
>>> nltk.download('punkt')

```

For more information see: <https://www.nltk.org/data.html>

Attempted to load tokenizers/punkt/PY3/english.pickle

Searched in:

- '/home/jovyan/nltk\_data'
- '/opt/conda/nltk\_data'
- '/opt/conda/share/nltk\_data'
- '/opt/conda/lib/nltk\_data'
- '/usr/share/nltk\_data'

```
- '/usr/local/share/nltk_data'
- '/usr/lib/nltk_data'
- '/usr/local/lib/nltk_data'
- ''
```

```
*****
```

```
In [12]: results=cross_validate(train_data, 10) # will work and output overall pe
results
```

```
Out[12]: []
```

## 4. Error Analysis (10 marks)

```
In [13]: from sklearn import metrics
import matplotlib.pyplot as plt
# a function to make the confusion matrix readable and pretty
def confusion_matrix_heatmap(y_test, preds, labels):
    """Function to plot a confusion matrix"""
    # pass labels to the confusion matrix function to ensure right order
    cm = metrics.confusion_matrix(y_test, preds, labels=labels)
    fig = plt.figure(figsize=(10,10))
    ax = fig.add_subplot(111)
    cax = ax.matshow(cm)
    plt.title('Confusion matrix of the classifier')
    fig.colorbar(cax)
    ax.set_xticks(np.arange(len(labels)))
    ax.set_yticks(np.arange(len(labels)))
    ax.set_xticklabels( labels, rotation=45)
    ax.set_yticklabels( labels)

    for i in range(len(cm)):
        for j in range(len(cm)):
            text = ax.text(j, i, cm[i, j],
                           ha="center", va="center", color="w")

    plt.xlabel('Predicted')
    plt.ylabel('True')

    # fix for mpl bug that cuts off top/bottom of seaborn viz:
    b, t = plt.ylim() # discover the values for bottom and top
    b += 0.5 # Add 0.5 to the bottom
    t -= 0.5 # Subtract 0.5 from the top
    plt.ylim(b, t) # update the ylim(bottom, top) values
    plt.show() # ta-da!
    plt.show()
```

```
In [14]: preds=[]
for i in range(len(test_data)):
    pred=predict_labels(test_data[i][0],a)
    preds.append(pred)
y_test=np.array(test_data)[: ,1]
y_test
preds
```

```
-----
---
IndexError                                Traceback (most recent call la
st)
Cell In [14], line 5
      3     pred=predict_labels(test_data[i][0],a)
      4     preds.append(pred)
----> 5 y_test=np.array(test_data)[: ,1]
      6 y_test
      7 preds

IndexError: too many indices for array: array is 1-dimensional, but 2 we
re indexed
```

```
In [15]: confusion_matrix_heatmap(preds,y_test,['REAL','FAKE'])
```

```
-----
---
NameError                                Traceback (most recent call la
st)
Cell In [15], line 1
----> 1 confusion_matrix_heatmap(preds,y_test,['REAL','FAKE'])

NameError: name 'y_test' is not defined
```

## Questions 5 (20%) and 6 (20%) (recommend starting a new notebook)

```
In [16]: # Finally, check the accuracy of your classifier by training on all the t
# and testing on the test set
# Will only work once all functions are complete
functions_complete = True # set to True once you're happy with your meth
if functions_complete:
    print(test_data[0]) # have a look at the first test data instance
    classifier = train_classifier(train_data) # train the classifier
    test_true = [t[1] for t in test_data] # get the ground-truth labels
    test_pred = predict_labels([x[0] for x in test_data], classifier) #
    final_scores = precision_recall_fscore_support(test_true, test_pred,
    print("Done training!")
    print("Precision: %f\nRecall: %f\nF Score:%f" % final_scores[:3])
```

```
-----  
-----  
IndexError                                Traceback (most recent call la  
st)  
Cell In [16], line 6  
      4 functions_complete = True # set to True once you're happy with  
your methods for cross val  
      5 if functions_complete:  
----> 6     print(test_data[0]) # have a look at the first test data i  
nstance  
      7     classifier = train_classifier(train_data) # train the class  
ifier  
      8     test_true = [t[1] for t in test_data] # get the ground-tru  
th labels from the data  
  
IndexError: list index out of range
```

In [ ]: